

SelfDrivingCar

February 12, 2019

```
In [1]: import numpy as np
import pandas as pd
import os
import matplotlib.pyplot as plt
import seaborn as sns
import imageio
import scipy.misc
from tqdm import tqdm
import warnings
from tensorflow.core.protobuf import saver_pb2
from tensorflow.python.client import device_lib

In [2]: warnings.filterwarnings('ignore', category=FutureWarning)

In [3]: !gpustat

instance-gpu  Mon Feb 11 16:12:23 2019
[0] Tesla P100-PCIE-16GB | 34'C,  0 % |    82 / 16280 MB | root(82M)

In [15]: !cat ./dataset2/data.txt | wc -l

45406

In [16]: !head -n 1 ./dataset2/data.txt

0.jpg 0.000000

In [17]: steer_df = pd.read_csv('./dataset2/data.txt', sep=' ',
                                names=['image_file', 'steer_degree'],
                                engine='python')

In [18]: steer_df['image_file'] = './dataset2/' + steer_df.image_file

In [96]: steer_df.shape

Out[96]: (45406, 3)
```

```
In [19]: steer_df.head()
```

```
Out[19]:
```

	image_file	steer_degree
0	./dataset2/0.jpg	0.0
1	./dataset2/1.jpg	0.0
2	./dataset2/2.jpg	0.0
3	./dataset2/3.jpg	0.0
4	./dataset2/4.jpg	0.0

Convert degrees to Radians

```
In [21]: steer_df['steer_rad'] = steer_df.steer_degree * 0.01745
```

```
In [22]: steer_df.steer_rad.unique()
```

```
Out[22]: array([ 0.          ,  0.204165 ,  0.1813055, ..., -1.4781895, -1.372617 ,
                -0.883319 ])
```

```
In [23]: steer_df.shape
```

```
Out[23]: (45406, 3)
```

```
In [27]: train = steer_df.iloc[0:31785, :] # More than 31785 for epoch calculations
        test = steer_df.iloc[31785:, :]
```

```
In [28]: train.shape
```

```
Out[28]: (31785, 3)
```

```
In [29]: test.shape
```

```
Out[29]: (13621, 3)
```

0.0.1 Convert image into feature matrix

```
In [30]: # def _convert_image(image_path, strip, resize, normalize):
        #     img_matrix = imageio.imread(image_path)
        #     if strip is not None:
        #         img_matrix = img_matrix[-strip:]
        #     if resize is not None:
        #         img_matrix = scipy.misc.imresize(img_matrix, resize)
        #     if not normalize:
        #         return img_matrix
        #     return img_matrix / 255.0
```

```
In [106]: # train_batch_pointer = 0
        # test_batch_pointer = 0
```

```

In [133]: # def get_batch_data(size, subset='train', **kwargs):
#         global train_batch_pointer
#         global test_batch_pointer
#         strip = kwargs.get('strip')
#         resize = kwargs.get('resize')
#         normalize = kwargs.get('normalize', True)
#         train_pointer = kwargs.get('train_pointer', -1)
#         test_pointer = kwargs.get('test_pointer', -1)

#         fetcher = kwargs.get('fetcher', 'sequential')
#         if fetcher not in ['sequential', 'randomized']:
#             raise NotImplementedError('Fetcher type is not implemented')

#         x = []

#         if subset == 'train':
#             if fetcher == 'randomized':
#                 train_index = np.random.randint(low=0, high=train.shape[0], size=size)
#             else:
#                 if train_batch_pointer == - 1 or train_batch_pointer > train.shape[0]:
#                     raise IndexError('Train pointer value is less/more than it should be')
#                 train_index = np.arange(train_batch_pointer, train_batch_pointer + size)
#                 print(train_pointer, train_pointer+size)
#                 images = train.image_file.values[train_index]
#                 y = train.steer_rad.values[train_index]
#                 train_batch_pointer += size

#         elif subset == 'test':
#             if fetcher == 'randomized':
#                 train_index = np.random.randint(low=0, high=test.shape[0], size=size)
#             else:
#                 if test_batch_pointer == -1 or test_batch_pointer > test.shape[0]:
#                     raise IndexError('Test pointer value is less/more than it should be')
#                 test_index = np.arange(test_batch_pointer, test_batch_pointer + size)
#                 images = test.image_file.values[test_index]
#                 y = test.steer_rad.values[test_index]
#                 test_batch_pointer += size
#         for image in images:
#             x.append(_convert_image(image, strip, resize, normalize))

#         return np.array(x), y.reshape(-1, 1)
import random
xs = []
ys = []

train_batch_pointer = 0
val_batch_pointer = 0
#read data.txt

```

```

with open("./dataset2/data.txt") as f:
    for line in f:
        xs.append("./dataset2/" + line.split()[0])
        #the paper by Nvidia uses the inverse of the turning radius,
        #but steering wheel angle is proportional to the inverse of turning radius
        #so the steering wheel angle in radians is used as the output
        ys.append(float(line.split()[1]) * scipy.pi / 180)

#get number of images
num_images = len(xs)

#shuffle list of images
c = list(zip(xs, ys))
random.shuffle(c)
xs, ys = zip(*c)

train_xs = xs[:int(len(xs) * 0.8)]
train_ys = ys[:int(len(xs) * 0.8)]

val_xs = xs[-int(len(xs) * 0.2):]
val_ys = ys[-int(len(xs) * 0.2):]

num_train_images = len(train_xs)
num_val_images = len(val_xs)

def LoadTrainBatch(batch_size):
    global train_batch_pointer
    x_out = []
    y_out = []
    for i in range(0, batch_size):
        x_out.append(scipy.misc.imresize(
            scipy.misc.imread(train_xs[(
                train_batch_pointer + i) % num_train_images])[-150:], [66, 200]) / 255
            y_out.append([train_ys[(train_batch_pointer + i) % num_train_images]])
        train_batch_pointer += batch_size
    return x_out, y_out

def LoadValBatch(batch_size):
    global val_batch_pointer
    x_out = []
    y_out = []
    for i in range(0, batch_size):
        x_out.append(scipy.misc.imresize(
            scipy.misc.imread(
                val_xs[(val_batch_pointer + i) % num_val_images])[-150:], [66, 200]) /
            y_out.append([val_ys[(val_batch_pointer + i) % num_val_images]])
        val_batch_pointer += batch_size

```

```
    return x_out, y_out
```

0.0.2 Build the Conv Net

```
In [134]: import tensorflow as tf
```

```
In [135]: def get_weights(shape, name, init='glorot_uniform'):
    """
    Get weights and initialize them
    for the each layer.
    """
    if init == 'glorot_uniform':
        i = tf.glorot_uniform_initializer(seed=42)
    elif init == 'glorot_normal':
        i = tf.glorot_normal_initializer(seed=42)
    elif init == 'xavier_uniform':
        i = tf.contrib.layers.xavier_initializer(seed=42)
    elif init == 'xavier_normal':
        i = tf.contrib.layers.xavier_initializer(seed=42, uniform=False)
    elif init == 'he_uniform':
        i = tf.keras.initializers.he_uniform(seed=42)
    elif init == 'he_normal':
        i = tf.keras.initializers.he_normal(seed=42)
    elif init == 'raw':
        if len(shape) > 1:
            initial = tf.truncated_normal(shape, stddev=0.1)
            return tf.Variable(initial)
        else:
            initial = tf.constant(0.1, shape=shape)
            return tf.Variable(initial)
    # i = tf.initializers.random_normal(seed=42)

    return tf.get_variable(
        name=name,
        shape=shape,
        dtype=tf.float32,
        initializer=i,)

def flatten(X, size):
    return tf.reshape(X, [-1, size])

def Dense(X, size, init, name, activation):
    w = get_weights(shape=size, name='W_' + name, init=init)
    b = get_weights(shape=[size[-1]], name='b_' + name, init=init)

    dense = tf.matmul(X, w) + b
    print(name, size, size[-1])
```

```

    ## Applying activation

    if activation == 'relu':
        h_fc = tf.nn.relu(dense)
    elif activation == 'sigmoid':
        h_fc = tf.nn.sigmoid(dense)
    elif activation == 'leaky_relu':
        h_fc = tf.nn.leaky_relu(dense)
    elif activation == 'tanh':
        h_fc = tf.nn.tanh(dense)
    elif activation == 'atan':
        h_fc = tf.atan(dense)

    #     if dropout >= 0.0 and dropout < 1.0:
    #         return tf.nn.dropout(h_fc, keep_prob=dropout)
    return h_fc

def Conv2d(X, size, stride, init, name, padding, activation):
    """
    Get a conv layer on X for weight W and bias b
    with stride and padding
    """
    print(name, size, size[-1])
    w = get_weights(shape=size, name='W_' + name, init=init)
    b = get_weights(shape=[size[-1]], name='b_' + name, init=init)

    conv = tf.nn.conv2d(X, w, strides=[1, stride, stride, 1],
                        padding=padding) + b

    ## Applying activation

    if activation == 'relu':
        h_conv = tf.nn.relu(conv)
    elif activation == 'sigmoid':
        h_conv = tf.nn.sigmoid(conv)
    elif activation == 'leaky_relu':
        h_conv = tf.nn.leaky_relu(conv)

    return h_conv

```

```
In [136]: # get_weights((5, 5, 3, 24), name='conv2d_1').shape
```

```
In [137]: # get_weights((24), name='bias_conv2d_1').shape
```

```
In [138]: def get_available_gpus():
    local_device_protos = device_lib.list_local_devices()
    return [x.name for x in local_device_protos if x.device_type == 'GPU']

```

0.0.3 Network 1

```
In [139]: tf.reset_default_graph() # Reset any existing computation graphs

In [140]: x = tf.placeholder(name='input', dtype=tf.float32, shape=[None, 66, 200, 3])
          y_ = tf.placeholder(name='output', dtype=tf.float32, shape=[None, 1])
          keep_prob = tf.placeholder(tf.float32)

In [141]: x_image = x

In [142]: # first convolutional layer
          h_conv1 = Conv2d(x_image, (5, 5, 3, 24), 2,
                           init='raw', name='conv2d_1',
                           padding='VALID', activation='relu')

          # second convolutional layer
          h_conv2 = Conv2d(h_conv1, (5, 5, 24, 36), 2,
                           init='raw', name='conv2d_2',
                           padding='VALID', activation='relu')

          # third convolutional layer
          h_conv3 = Conv2d(h_conv2, (5, 5, 36, 48), 2,
                           init='raw', name='conv2d_3',
                           padding='VALID', activation='relu')

          # fourth convolutional layer
          h_conv4 = Conv2d(h_conv3, (3, 3, 48, 64), 1,
                           init='raw', name='conv2d_4',
                           padding='VALID', activation='relu')

          # fifth convolutional layer
          h_conv5 = Conv2d(h_conv4, (3, 3, 64, 64), 1,
                           init='raw', name='conv2d_5',
                           padding='VALID', activation='relu')

          # Flatten layer
          h_conv5_flatten = flatten(h_conv5, size=1152)

          # Dense layer 1
          h_fc1 = Dense(h_conv5_flatten, (1152, 1164), name='dense1',
                        init='raw',
                        activation='relu')

          # Dropout 1
          h_fc1_drop = tf.nn.dropout(h_fc1, keep_prob)

          # Dense Layer 2
          h_fc2 = Dense(h_fc1_drop, (1164, 100), name='dense2',
                        init='raw',
                        activation='relu')
```

```

# Dropout 2
h_fc2_drop = tf.nn.dropout(h_fc2, keep_prob)

# Dense Layer 3
h_fc3 = Dense(h_fc2_drop, (100, 50), name='dense3',
              init='raw',
              activation='relu')

# Dropout 3
h_fc3_drop = tf.nn.dropout(h_fc3, keep_prob)

# Dense Layer 4
h_fc4 = Dense(h_fc3_drop, (50, 10), name='dense4',
              init='raw',
              activation='relu')

# Dropout 4
h_fc4_drop = tf.nn.dropout(h_fc4, keep_prob)

# Output
Y = Dense(h_fc4_drop, (10, 1), name='output',
          init='raw',
          activation='atan')
y = tf.multiply(Y, 2)

conv2d_1 (5, 5, 3, 24) 24
conv2d_2 (5, 5, 24, 36) 36
conv2d_3 (5, 5, 36, 48) 48
conv2d_4 (3, 3, 48, 64) 64
conv2d_5 (3, 3, 64, 64) 64
dense1 (1152, 1164) 1164
dense2 (1164, 100) 100
dense3 (100, 50) 50
dense4 (50, 10) 10
output (10, 1) 1

```

0.0.4 Train the network

```
In [143]: LOGDIR = './save/'
```

```
In [144]: get_available_gpus()
```

```
Out[144]: ['/device:GPU:0']
```

```
In [ ]: with tf.Session() as sess:
        L2NormConst = 0.001
        train_vars = tf.trainable_variables()
        loss = tf.reduce_mean(tf.square(tf.subtract(y_, y))) + tf.add_n(
            [tf.nn.l2_loss(v) for v in train_vars]) * L2NormConst

```



```

train_step = tf.train.AdamOptimizer(1e-4).minimize(loss)
sess.run(tf.global_variables_initializer())

# create a summary to monitor cost tensor
tf.summary.scalar(name="RMSE", tensor=loss)
# merge all summaries into a single op
merged_summary_op = tf.summary.merge_all()
saver = tf.train.Saver(write_version = saver_pb2.SaverDef.V1)

logs_path = './logs'
summary_writer = tf.summary.FileWriter(logdir=logs_path, graph=tf.get_default_graph())
epochs = 30
batch_size = 100
with tf.device('/gpu:0'):
    for epoch in range(epochs):
        for i in range(int(num_images/batch_size)):
            xs, ys = LoadTrainBatch(batch_size)
            train_step.run(feed_dict={x: xs, y_: ys, keep_prob: 0.8})
            if i % 10 == 0:
                xs, ys = LoadValBatch(batch_size)
                loss_value = loss.eval(feed_dict={x: xs, y_: ys, keep_prob: 1.0})
                print("Epoch: %d, Step: %d, Loss: %g" % (epoch, epoch * batch_size + i, loss_value))

        # write logs at every iteration
        summary = merged_summary_op.eval(feed_dict={x:xs, y_: ys, keep_prob: 1.0})
        summary_writer.add_summary(summary, epoch * num_images/batch_size + i)

        if i % batch_size == 0:
            if not os.path.exists(LOGDIR):
                os.makedirs(LOGDIR)
            checkpoint_path = os.path.join(LOGDIR, "model.ckpt")
            filename = saver.save(sess, checkpoint_path)
            print("Model saved in file: %s" % filename)
print("Run the command line:\n" \
      "--> tensorboard --logdir=./logs " \
      "\nThen open http://0.0.0.0:6006/ into your web browser")

```

Last 5 steps of the last epoch

```

Epoch: 29, Step: 3310, Loss: 0.146305
Epoch: 29, Step: 3320, Loss: 0.137446
Epoch: 29, Step: 3330, Loss: 0.132673
Epoch: 29, Step: 3340, Loss: 0.130797
Epoch: 29, Step: 3350, Loss: 0.443287

```

0.0.5 Network 2

In [154]: `tf.reset_default_graph()` *# Reset any existing computation graphs*

```

In [155]: x = tf.placeholder(name='input', dtype=tf.float32, shape=[None, 66, 200, 3])
          y_ = tf.placeholder(name='output', dtype=tf.float32, shape=[None, 1])
          keep_prob = tf.placeholder(tf.float32)

In [156]: x_image = x

In [157]: # first convolutional layer
          h_conv1 = Conv2d(x_image, (5, 5, 3, 24), 2,
                           init='raw', name='conv2d_1',
                           padding='VALID', activation='relu')

          # second convolutional layer
          h_conv2 = Conv2d(h_conv1, (5, 5, 24, 36), 2,
                           init='raw', name='conv2d_2',
                           padding='VALID', activation='relu')

          # third convolutional layer
          h_conv3 = Conv2d(h_conv2, (5, 5, 36, 48), 2,
                           init='raw', name='conv2d_3',
                           padding='VALID', activation='relu')

          # fourth convolutional layer
          h_conv4 = Conv2d(h_conv3, (3, 3, 48, 64), 1,
                           init='raw', name='conv2d_4',
                           padding='VALID', activation='relu')

          # fifth convolutional layer
          h_conv5 = Conv2d(h_conv4, (3, 3, 64, 64), 1,
                           init='raw', name='conv2d_5',
                           padding='VALID', activation='relu')

          # Flatten layer
          h_conv5_flatten = flatten(h_conv5, size=1152)

          # Dense layer 1
          h_fc1 = Dense(h_conv5_flatten, (1152, 1164), name='dense1',
                        init='raw',
                        activation='relu')

          # Dropout 1
          h_fc1_drop = tf.nn.dropout(h_fc1, keep_prob)

          # Dense Layer 2
          h_fc2 = Dense(h_fc1_drop, (1164, 100), name='dense2',
                        init='raw',
                        activation='relu')

          # Dropout 2
          h_fc2_drop = tf.nn.dropout(h_fc2, keep_prob)

```

```

# Dense Layer 3
h_fc3 = Dense(h_fc2_drop, (100, 50), name='dense3',
              init='raw',
              activation='relu')

# Dropout 3
h_fc3_drop = tf.nn.dropout(h_fc3, keep_prob)

# Dense Layer 4
h_fc4 = Dense(h_fc3_drop, (50, 10), name='dense4',
              init='raw',
              activation='relu')

# Dropout 4
h_fc4_drop = tf.nn.dropout(h_fc4, keep_prob)

# Output
W_fc5 = get_weights([10, 1], name='W_fc5', init='raw')
b_fc5 = get_weights([1], name='b_fc5', init='raw')
y = tf.matmul(h_fc4_drop, W_fc5) + b_fc5

conv2d_1 (5, 5, 3, 24) 24
conv2d_2 (5, 5, 24, 36) 36
conv2d_3 (5, 5, 36, 48) 48
conv2d_4 (3, 3, 48, 64) 64
conv2d_5 (3, 3, 64, 64) 64
dense1 (1152, 1164) 1164
dense2 (1164, 100) 100
dense3 (100, 50) 50
dense4 (50, 10) 10

```

0.0.6 Train the network

```
In [158]: LOGDIR = './save/'
```

```
In [159]: get_available_gpus()
```

```
Out[159]: ['/device:GPU:0']
```

```

In [ ]: with tf.Session() as sess:
    L2NormConst = 0.001
    train_vars = tf.trainable_variables()
    loss = tf.reduce_mean(tf.square(tf.subtract(y_, y))) + tf.add_n(
        [tf.nn.l2_loss(v) for v in train_vars]) * L2NormConst
    train_step = tf.train.AdamOptimizer(1e-3).minimize(loss)
    sess.run(tf.global_variables_initializer())

    # create a summary to monitor cost tensor
    tf.summary.scalar(name="RMSE", tensor=loss)

```

```

# merge all summaries into a single op
merged_summary_op = tf.summary.merge_all()
saver = tf.train.Saver(write_version = saver_pb2.SaverDef.V1)

logs_path = './logs'
summary_writer = tf.summary.FileWriter(logdir=logs_path, graph=tf.get_default_graph())
epochs = 30
batch_size = 100
with tf.device('/gpu:0'):
    for epoch in range(epochs):
        for i in range(int(num_images/batch_size)):
            xs, ys = LoadTrainBatch(batch_size)
            train_step.run(feed_dict={x: xs, y_: ys, keep_prob: 0.5})
            if i % 10 == 0:
                xs, ys = LoadValBatch(batch_size)
                loss_value = loss.eval(feed_dict={x: xs, y_: ys, keep_prob: 1.0})
                print("Epoch: %d, Step: %d, Loss: %g" % (epoch, epoch * batch_size + i, loss_value))

# write logs at every iteration
summary = merged_summary_op.eval(feed_dict={x:xs, y_: ys, keep_prob: 1.0})
summary_writer.add_summary(summary, epoch * num_images/batch_size + i)

if i % batch_size == 0:
    if not os.path.exists(LOGDIR):
        os.makedirs(LOGDIR)
    checkpoint_path = os.path.join(LOGDIR, "model2.ckpt")
    filename = saver.save(sess, checkpoint_path)
    print("Model saved in file: %s" % filename)
print("Run the command line:\n" \
      "--> tensorboard --logdir=./logs " \
      "\nThen open http://0.0.0.0:6006/ into your web browser")

```

Last 5 steps of last epoch

```

Epoch: 29, Step: 3310, Loss: 0.0919582
Epoch: 29, Step: 3320, Loss: 0.136387
Epoch: 29, Step: 3330, Loss: 0.193896
Epoch: 29, Step: 3340, Loss: 0.134995
Epoch: 29, Step: 3350, Loss: 0.126513

```

For video link: <https://youtu.be/zhbPPFHbLDE>

0.0.7 Saved Models

In [8]: !ls -l save

```

total 50360
-rw-rw-r-- 1 mayukhpay mayukhpay          79 Feb 11 23:40 checkpoint
-rw-rw-r-- 1 mayukhpay mayukhpay 31920357 Feb 11 23:40 model2.ckpt

```

```
Autopilot-TensorFlow-master — IPython: mayukhsarkar/Downloads — -bash — 120x35
Steering angle: 12.89789187957842 (pred) 15.830000000000001 (actual)
Steering angle: 12.94597926429896 (pred) 15.83 (actual)
Steering angle: 12.859865669241483 (pred) 16.24 (actual)
Steering angle: 11.532397045614887 (pred) 16.64 (actual)
Steering angle: 11.21565560592008 (pred) 16.74 (actual)
Steering angle: 9.54359451592162 (pred) 16.84 (actual)
Steering angle: 9.355584167506912 (pred) 16.84 (actual)
Steering angle: 10.296709956827446 (pred) 16.84 (actual)
Steering angle: 11.184702552619736 (pred) 16.939999999999998 (actual)
Steering angle: 14.398301873684548 (pred) 17.04 (actual)
Steering angle: 14.51225148614054 (pred) 17.239999999999995 (actual)
Steering angle: 16.077918674109636 (pred) 17.55 (actual)
Steering angle: 13.761833969956871 (pred) 17.75 (actual)
Steering angle: 14.063504642576117 (pred) 17.850000000000005 (actual)
Steering angle: 14.117311165310184 (pred) 17.850000000000005 (actual)
Steering angle: 13.95713298398963 (pred) 17.850000000000005 (actual)
Steering angle: 13.35778217677372 (pred) 17.850000000000005 (actual)
Steering angle: 14.740687162931616 (pred) 17.850000000000005 (actual)
Steering angle: 14.847656463070479 (pred) 17.850000000000005 (actual)
Steering angle: 15.79662245578454 (pred) 17.850000000000005 (actual)
Steering angle: 17.9554796203382 (pred) 17.850000000000005 (actual)
Steering angle: 18.7515569511234 (pred) 17.850000000000005 (actual)
Steering angle: 18.508214386468772 (pred) 17.34 (actual)
Steering angle: 16.993934126720266 (pred) 17.04 (actual)
Steering angle: 18.080955318033208 (pred) 16.64 (actual)
Steering angle: 17.93729424167304 (pred) 16.13 (actual)
Steering angle: 16.775873507278437 (pred) 15.929999999999998 (actual)
Steering angle: 15.52632796466504 (pred) 15.929999999999998 (actual)
Steering angle: 15.153507211461772 (pred) 15.929999999999998 (actual)
Steering angle: 13.092547148281907 (pred) 15.929999999999998 (actual)
Steering angle: 12.308332148537723 (pred) 15.929999999999998 (actual)
Steering angle: 10.927679417258638 (pred) 15.83 (actual)
Steering angle: 9.924468062156334 (pred) 16.03 (actual)
Steering angle: 10.204516063410614 (pred) 16.03 (actual)
Steering angle: 9.557429064085483 (pred) 16.03 (actual)
```

Output

```
-rw-rw-r-- 1 mayukhpay mayukhpay 317497 Feb 11 23:40 model2.ckpt.meta
-rw-rw-r-- 1 mayukhpay mayukhpay 19152108 Feb 11 21:13 model.ckpt
-rw-rw-r-- 1 mayukhpay mayukhpay 166715 Feb 11 21:13 model.ckpt.meta
```

0.0.8 Observations

- The initial model works better
- The updated model is performing well but not as good as the initial one.
- I hardly think not doing atan at output was the main reason.
- Because now we are directly taking values instead of atan and square of it, the loss on paper is better than the previous one because I am getting more granular control over the output instead of a smoothed out value.
- I guess the lower performance is due to the learning rate and dropout rate changes.

Below you can see the sample output.